

# Introduction to Acceleration Structures


## Contents

[Introduction](#)[Bounding Volume](#)[Bounding Volume Hierarchy: BVH \(part 1\)](#)[Bounding Volume Hierarchy: BVH \(part 2\)](#)[Grid](#)[What Else?](#)[Source Code](#)

The rest of Kay and Kajiya's paper is focused on improving their technique by grouping the bounding volumes we have described in the previous chapter into a hierarchy of bounding volumes. This technique is called a **Bounding Volume Hierarchy** or **BVH**. It generally provides (compared to other possible acceleration structures) very good results. Many variations based on this principle exist. In this chapter we will look at the method used and described by Kay and Kajiya.

The idea of grouping bounding volumes into larger volumes which we themselves group, etc. is very simple and also quite easy to understand. In figure 2, we are showing a group of objects associated with their respective bounding volumes (here a box for simplicity). By merging these bounding boxes together (usually by proximity) we obtain larger bounding volumes representing groups of objects. It is easy to see that if a ray doesn't intersect any of these larger groups, then we can avoid testing the volumes enclosed by these groups, possibly rejecting many objects at once. This obviously, saves a lot of render time.

The principle is very simple however to be efficient, the objects and the volumes have to be grouped by proximity. The problem if they



© [www.scratchapixel.com](http://www.scratchapixel.com)

Figure 1: grouping bounding volumes.

are not grouped by proximity, is illustrated in figure 3. The two bounding boxes of the red teapots have been grouped together even though they are far away from each other. In figure 3, the ray intersects the two bounding volumes and four teapots have to be tested for an intersection with the ray. In figure 1, only two of these teapot are tested for an intersection (D and E).

To group objects by proximity, Kay and Kajiya propose to insert them in a **space partitioning data structure**. This structure partitions space in sub-regions and objects are inserted in these sub-regions usually based on their position. The process is illustrated in figure 3. If we create a box which encloses all the objects of the scene then we can subdivide this box in eight equally sized sub-boxes (in figures 4 and 5, this process is illustrated in 2D). Each object of the scene can then be inserted (in the order they have been added to the scene) in the sub-boxes they overlap. We can then assume that all the objects contained in one of the sub-voxels are pretty close to each other (at least closer than the objects which are in the other sub-boxes).

As you can see in figure 4, an object may overlap several sub-boxes or cells. In that case, Kay and Kajiya arbitrarily insert the object in one of the overlapped cells. We have chosen to insert them in the cell in which lies the object's bounding box centroid (represented by the dot in the center of each teapot). Two spatial data structures are proposed in the paper: a **median cut** and an **octree**. The first is related somehow to the k-d tree space partitioning data structure which you may have heard about. We won't be presenting these structures in this lesson. The latter, the octree, will be used in our implementation of Kay's and Kajiya's paper. As described before, it partitions a cube into eight cells which are themselves subdivided etc. The recursion process stops when all the objects have been






Figure 2: example of a bounding volume hierarchy. Objects contained in bounding box C are don't need to be tested.



© www.scratchapixel.com

Figure 3: space is subdivided into smaller regions.



© www.scratchapixel.com

Figure 4: an octree in two dimensions is a quadtree. The node is divided into four cells, and objects are inserted in the cell they overlap. The process of subdividing the cells can be repeated as many times as we want. For example until there is one object

inserted in the octree or when we reach a maximum user-defined depth.

per cell or when we have reached a maximum user-defined depth.

## Building the Hierarchy

The octree data structure is presented in detail in an other lesson. However, in short, the process for the creation and insertion of the object is the following. First we compute the bounding volumes of all the objects in the scene and as we go, we compute the overall scene bounding volume which is the result of these volumes combined. From that scene overall volume we can define the size of the octree (a cube located at the center of the overall volume, whose dimension is the maximum value of any of the volume's extent along the x- y- and z-axis). This cube represents the top node of the octree, its **root**. When we then insert objects in this octree, we traverse the tree in a **top-down** fashion. If the node in which we want to insert an object is a leaf, and that the leaf doesn't contain any object yet, then we insert the object in this **node**. If the node already contains an object (and unless we have reached a maximum user defined depth) then the current node is split in eight cells and the object contained by the node as well as the new object are inserted in the cells (note the object may overlap more than one cell) they overlap.

This process is repeated until all the objects are inserted in the octree. In a second step, the octree is traversed in a **bottom-up** fashion this time. We start from the leaves, and compute the overall bounding box of the objects they contain (a leaf can have one or more objects). The overall volume box of the node above the leaves is then computed from combining its children bounding volumes. This process is repeated until we reach the root (which at this point should be the same as the extent of the overall scene we have computed earlier on). At the end of these two processes, we obtain a representation of the scene as a **hierarchy of bounding volumes**, where the volumes are grouped based on proximity (objects contained in the leaves of a node are grouped together, etc.). Note that the octree is actually not used as an acceleration structure (do not confuse this technique with the octree used as an acceleration structure - we will write a lesson on this technique in the future). The octree is actually of no real benefit for the ray-intersection tests. It is only used to compute and build this hierarchy of volumes.



Figure 5: objects inserted in a quadtree.

## Intersecting the Hierarchy

The intersection of the BVH is quite simple. We start from the root node and test if the ray intersects the bounding volume for this node.

Generally, if the bounding volume of a node is intersected by the ray, we go one level down in the tree and test for an intersection with the bounding volumes of this node's children. If the node is a leaf, we test if the ray intersects any of the objects it contains. This solution is quite simple to implement however, it can be further optimized. When we test for the intersection of a ray with the bounding volumes of a node's children, assuming that the bounding volumes of more than one of these cells have been intersected, we should first investigate the children of the node with the smallest intersection distance, as the visible object is most likely to be contained in these cells. This idea is illustrated in figure 6 (we have represented the intersection with the bounding volume of a cell by a point at the intersection of the ray with the cell's boundaries. Keep in mind that in reality, we have an intersection with the bounding volume of the cell, represented by spheres in the figure, not the cell itself). As you can see, the ray intersects the root node (in black) which leads to testing the root's children (in red). The cells are tested in the order they have been created: 0, 1, 2, 3. The ray intersects the bounding volumes of cell 0 and 1 and therefore the children of these nodes are tested next. However the problem is that the children of the cells 0 will be tested before the children of cell 1 even though the intersection distance with the bounding volume from cell 1 ( $t_1$ ) is smaller than the intersection distance with the volume from cell 0 ( $t_0$ ). It would be more efficient to first test the children from cell 1 before testing the children from cell 0 as it would lead to finding the intersection with the teapot A sooner. Kay and Kajiya propose to use a list in which the node's children which are pierced by the ray are inserted and sorted according to their intersection distance (from the smallest to the largest). We then remove the first node from this list and test its children. If any of these children are intersected by the ray, they get themselves inserted to the list. They will be inserted in




Figure 6: intersecting with a quadtree. If a cell is intersected the cell's children are tested for an intersection. This process continues until a leaf node is intersected. The geometry contained in this leaf node is then tested.

front of the existing elements from the list since their intersection distance is necessarily smaller than the intersection distance of any of the nodes already inserted.

In programming such lists are called **priority queues**. They are like regular queues where each element has a "priority" associated with it. In a priority queue, an element with high priority is served before an element with low priority. If two elements have the same priority, they are served according to their order in the queue. In our particular case, we will use the intersection distance to the bounding volume to define the priority of the volumes in the queue: the volume with the smallest intersection distance is the element from the queue with the highest priority. In C++ we can use the `priority_queue` class which is part of the STL C++ standard (see the source code section for more details).

By following this process, we are sure to always test the nodes from the hierarchy which are the closest to the ray's origin first. If these nodes do not contain the visible object, we then keep going by testing the nodes which are further away from the ray's origin, etc. until the list is empty.

```

001 while (priority_queue is not empty) {
002     Node node = get node from the priority_queue with highest priority (and remove it)
003     if node is a leaf
004         for each object contained in node
005             if object is intersected by ray
006                 keep this object as potential visible object
007     if an object was intersected
008         return
009     else
010         // keep traversing the hierarchy by looking down the node's children
011         for each children in node
012             if node's children is intersected by ray at distance t
013                 insert node into priority_queue (use intersection distance t to define
014 }

```

There is one last detail we need to be careful about. If you look at figure 7, you can see that we can't use the intersection distance to the bounding volumes to decide for certain which ones of these volumes contain the visible object. Even though the distance to the bounding volume of B  $t_{VB}$  is lower than the distance to the bounding volume of C  $t_{VC}$ , the object that the ray will intersect is C and not B. We only know for certain that we can stop the traversal of the hierarchy when the distance to the intersected object, is lower than the intersection distance to the bounding volume of the next node in the list. As illustrated in figure 7, the nodes in the priority list should be




Figure 7: the first intersected bounding volume doesn't necessarily contain the nearest intersected object.

in the following order VB, VC, VA (because  $t_{VB} < t_{VC} < t_{VA}$ ). When we intersect the object contained by VB (the bounding volume for B) we find the intersection distance  $t_B$  for the object B. However  $t_B$  is not smaller than  $t_{VC}$  therefore we also have to test for the intersection with the object C and we find that  $t_C$  is lower than  $t_B$  therefore C becomes the potential intersected object instead of B. But because  $t_C$  is lower than  $t_{VA}$  we don't need to test for an intersection with A nor do we need to test for an intersection with any of the nodes which might be in the list after VA. We can therefore stop the hierarchical traversal and return B as the visible object.

If we use all these techniques in our ray tracer we now get the following results:

Render time	: 1.61 (sec)
Total number of triangles	: 16384
Total number of primary rays	: 307200
Total number of ray-triangles tests	: 41341952
Total number of ray-triangles intersections	: 59017
Total number of ray-volume tests	: 1531064

The render is now 1.8 times faster (compared to the render from the previous chapter). The number of ray-volumes tests is reduced by 6.4 (proving that the bounding volume hierarchy is very efficient), the number of ray-triangle tests by 1.95 and the number of ray-triangle intersection by 1.89.

## Source Code

The complete source code is available in the last chapter of this lesson.

The BVH class has become more complex. We have added to the class an `OctreeNode` and an `Octree` class. Readers interested in learning about octaves are referred to following sections where they can find a lesson on this topic. The octree node holds eight pointers to other octree nodes which are its children. The class constructor takes the scene extent as an input parameter (line 38). This extent is used to compute and set the root node's dimension and position (its centroid. Lines 40-49). New objects (more precisely their bounding volume, but remember that the `Extents` class holds a pointer to the enclosed object) are inserted from the root node (lines 52-53). If the current node is an empty leaf (no objects had been inserted in this leaf yet), we insert the object in this leaf and return. If the node is a leaf and contains at least one or more object(s) but that we have reached the octree maximum depth, we nevertheless insert the object in this node (line 73). However, if we haven't reached the maximum depth yet, we tag the node as "internal" and re-insert the object already held by the node as well as the new object in the octree (lines 77-82). If the node is not a leaf, we check in which children from the node the object should be inserted in (lines 86-93). If the child doesn't exist yet we create it (lines 97-98) and finally we insert the object into this child node (line 99).

The function for building the hierarchy of volumes is simple. We start from the root and go down the tree until we reach the leaf nodes. The bounding volume of this node is computed by combining the bounding volumes of all the objects it keeps a pointer to.

Then, as we move back up again, the bounding volume of the current node is computed by combining the bounding volumes of its children (line 130).

```

001  class BVH : public AccelerationStructure
002  {
003      static const uint8_t kNumPlaneSetNormals = 7;
004      static const Vec3f planeSetNormals[kNumPlaneSetNormals];
005      struct Extents
006      {
007          Extents()
008          {
009              for (uint8_t i = 0; i < kNumPlaneSetNormals; ++i)
010                  d[i][0] = kInfinity, d[i][1] = -kInfinity;
011          }
012          void extendBy(const Extents &extents)
013          {
014              for (uint8_t i = 0; i < kNumPlaneSetNormals; ++i) {
015                  if (extents.d[i][0] < d[i][0]) d[i][0] = extents.d[i][0];
016                  if (extents.d[i][1] > d[i][1]) d[i][1] = extents.d[i][1];
017              }
018          }
019          bool intersect(
020              const float *precomputedNumerator, const float *precomputeDenominator,
021              float &tNear, float &tFar, uint8_t &planeIndex);
022          float d[kNumPlaneSetNormals][2]; // d values for each plane-set normals
023          const Object *object; // pointer contained by the volume (used by octree)
024      };
025      Extents *extents;
026      struct OctreeNode
027      {
028          OctreeNode *child[8];
029          std::vector<const Extents *> data;
030          Extents extents;
031          bool isLeaf;
032          uint8_t depth; // just for debugging
033          OctreeNode() : isLeaf(true) { memset(child, 0x0, sizeof(OctreeNode *) * 8); }
034          ~OctreeNode() { for (uint8_t i = 0; i < 8; ++i) if (child[i] != NULL) delete child[i]; }
035      };
036      struct Octree
037      {
038          Octree(const Extents &extents) : root(NULL)
039          {
040              float xdiff = extents.d[0][1] - extents.d[0][0];
041              float ydiff = extents.d[1][1] - extents.d[1][0];
042              float zdiff = extents.d[2][1] - extents.d[2][0];
043              float dim = std::max(xdiff, std::max(ydiff, zdiff));
044              Vec3f centroid(
045                  (extents.d[0][0] + extents.d[0][1]),
046                  (extents.d[1][0] + extents.d[1][1]),
047                  (extents.d[2][0] + extents.d[2][1]));
048              bounds[0] = (Vec3f(centroid) - Vec3f(dim)) * 0.5f;
049              bounds[1] = (Vec3f(centroid) + Vec3f(dim)) * 0.5f;
050              root = new OctreeNode;
051              root->extents = extents;
052              root->depth = 0;
053              root->isLeaf = false;
054              root->data.clear();
055              root->data.push_back(extents);
056              root->child[0] = new OctreeNode;
057              root->child[1] = new OctreeNode;
058              root->child[2] = new OctreeNode;
059              root->child[3] = new OctreeNode;
060              root->child[4] = new OctreeNode;
061              root->child[5] = new OctreeNode;
062              root->child[6] = new OctreeNode;
063              root->child[7] = new OctreeNode;
064              root->child[0]->data.push_back(extents);
065              root->child[1]->data.push_back(extents);
066              root->child[2]->data.push_back(extents);
067              root->child[3]->data.push_back(extents);
068              root->child[4]->data.push_back(extents);
069              root->child[5]->data.push_back(extents);
070              root->child[6]->data.push_back(extents);
071              root->child[7]->data.push_back(extents);
072              root->child[0]->depth = 1;
073              root->child[1]->depth = 1;
074              root->child[2]->depth = 1;
075              root->child[3]->depth = 1;
076              root->child[4]->depth = 1;
077              root->child[5]->depth = 1;
078              root->child[6]->depth = 1;
079              root->child[7]->depth = 1;
080              root->isLeaf = false;
081          }
082      };
083      Octree octree(extents);
084      OctreeNode *root = octree.root;
085      root->data.push_back(extents);
086      root->depth = 0;
087      root->isLeaf = false;
088      root->data.clear();
089      root->data.push_back(extents);
090      root->child[0] = new OctreeNode;
091      root->child[1] = new OctreeNode;
092      root->child[2] = new OctreeNode;
093      root->child[3] = new OctreeNode;
094      root->child[4] = new OctreeNode;
095      root->child[5] = new OctreeNode;
096      root->child[6] = new OctreeNode;
097      root->child[7] = new OctreeNode;
098      root->child[0]->data.push_back(extents);
099      root->child[1]->data.push_back(extents);
100      root->child[2]->data.push_back(extents);
101      root->child[3]->data.push_back(extents);
102      root->child[4]->data.push_back(extents);
103      root->child[5]->data.push_back(extents);
104      root->child[6]->data.push_back(extents);
105      root->child[7]->data.push_back(extents);
106      root->child[0]->depth = 1;
107      root->child[1]->depth = 1;
108      root->child[2]->depth = 1;
109      root->child[3]->depth = 1;
110      root->child[4]->depth = 1;
111      root->child[5]->depth = 1;
112      root->child[6]->depth = 1;
113      root->child[7]->depth = 1;
114      root->isLeaf = false;
115  }

```

```

050         root = new OctreeNode;
051     }
052     void insert(const Extents *extents)
053     { insert(root, extents, bounds[0], bounds[1], 0); }
054     void build()
055     { build(root, bounds[0], bounds[1]); }
056     ~Octree() { delete root; }
057     struct QueueElement
058     {
059         const OctreeNode *node; // octree node held by this node in the tree
060         float t; // used as key
061         QueueElement(const OctreeNode *n, float thit) : node(n), t(thit) {}
062         // comparator is > instead of < so priority_queue behaves like a min-heap
063         friend bool operator < (const QueueElement &a, const QueueElement &b) { ret
064     };
065     Vec3f bounds[2];
066     OctreeNode *root;
067 private:
068     void insert(
069         OctreeNode *node, const Extents *extents,
070         Vec3f boundMin, Vec3f boundMax, int depth)
071     {
072         if (node->isLeaf) {
073             if (node->data.size() == 0 || depth == 16) {
074                 node->data.push_back(extents);
075             }
076             else {
077                 node->isLeaf = false;
078                 while (node->data.size()) {
079                     insert(node, node->data.back(), boundMin, boundMax, depth);
080                     node->data.pop_back();
081                 }
082                 insert(node, extents, boundMin, boundMax, depth);
083             }
084         } else {
085             // insert bounding volume in the right octree cell
086             Vec3f extentsCentroid = (
087                 Vec3f(extents->d[0][0], extents->d[1][0], extents->d[2][0]) +
088                 Vec3f(extents->d[0][1], extents->d[1][1], extents->d[2][1])) * 0.5;
089             Vec3f nodeCentroid = (boundMax + boundMin) * 0.5f;
090             uint8_t childIndex = 0;
091             if (extentsCentroid[0] > nodeCentroid[0]) childIndex += 4;
092             if (extentsCentroid[1] > nodeCentroid[1]) childIndex += 2;
093             if (extentsCentroid[2] > nodeCentroid[2]) childIndex += 1;
094             Vec3f childBoundMin, childBoundMax;
095             Vec3f boundCentroid = (boundMin + boundMax) * 0.5;
096             computeChildBound(childIndex, boundCentroid, boundMin, boundMax, childE
097             if (node->child[childIndex] == NULL)
098                 node->child[childIndex] = new OctreeNode, node->child[childIndex]
099             insert(node->child[childIndex], extents, childBoundMin, childBoundMax,
100         }
101     }
102     void computeChildBound(

```



```

102         const uint8_t &i, const Vec3f &boundCentroid,
103         const Vec3f &boundMin, const Vec3f &boundMax,
104         Vec3f &pMin, Vec3f &pMax) const
105     {
106         pMin[0] = (i & 4) ? boundCentroid[0] : boundMin[0];
107         pMax[0] = (i & 4) ? boundMax[0] : boundCentroid[0];
108         pMin[1] = (i & 2) ? boundCentroid[1] : boundMin[1];
109         pMax[1] = (i & 2) ? boundMax[1] : boundCentroid[1];
110         pMin[2] = (i & 1) ? boundCentroid[2] : boundMin[2];
111         pMax[2] = (i & 1) ? boundMax[2] : boundCentroid[2];
112     }
113     // bottom-up construction
114     void build(OctreeNode *node, const Vec3f &boundMin, const Vec3f &boundMax)
115     {
116         if (node->isLeaf) {
117             // compute leaf node bounding volume
118             for (uint32_t i = 0; i < node->data.size(); ++i) {
119                 node->extents.extendBy(*node->data[i]);
120             }
121         }
122         else {
123             for (uint8_t i = 0; i < 8; ++i)
124                 if (node->child[i]) {
125                     Vec3f childBoundMin, childBoundMax;
126                     Vec3f boundCentroid = (boundMin + boundMax) * 0.5;
127                     computeChildBound(i, boundCentroid, boundMin, boundMax, childBc
128                     build(node->child[i], childBoundMin, childBoundMax);
129                     node->extents.extendBy(node->child[i]->extents);
130                 }
131             }
132         }
133     }
134 };
135 Octree *octree;
136 public:
137     BVH(const RenderContext *rcx);
138     const Object* intersect(const Ray<float> &ray, IsectData &isectData) const;
139     ~BVH();
140 };
141

```

In the intersect method of the BVH class (line 37) we first test if the ray intersects the bounding volume of the entire scene (the extent of the ochre's root node). If it does we initialise a `priority_queue` list with this node and the intersection distance from the ray's origin to its bounding volume. We overloaded the operator `<` (line 63 above) to make it behave like a **min heap** (by default it behaves like a max heap setting the element with the highest key value with the highest priority. We want the opposite). The rest of the code is similar to the pseudocode given above. We take the first node on top of the list (which we also remove from the list. Lines 72-73). If this node is a leaf then we test if

the ray intersects any of the objects contained by the node and we keep track of the intersection minimum distance (line 79). If it is an internal node, we test if the ray intersect the bounding volumes of the node's children, and when it does, we add the child node to the list (using the intersection distance as the key to set the element's priority in the list. Line 93). This process is repeated until the list is empty but can be stopped sooner if the intersection distance of the list's first node is greater than the distance to the intersected object (line 71).

```

001 BVH::BVH(const RenderContext *rcx) : AccelerationStructure(rcx), extents(NULL), octree(
002 {
003     Extents sceneExtents;
004     extents = new Extents[rcx->objects.size()];
005     for (uint32_t i = 0; i < rcx->objects.size(); ++i) {
006         for (uint8_t j = 0; j < kNumPlaneSetNormals; ++j) {
007             rcx->objects[i]->computeBounds(planeSetNormals[j], extents[i].d[j][0], exte
008         }
009         extents[i].object = rcx->objects[i];
010         sceneExtents.extendBy(extents[i]);
011     }
012     // create hierarchy
013     octree = new Octree(sceneExtents);
014     for (uint32_t i = 0; i < rcx->objects.size(); ++i) {
015         octree->insert(extents + i);
016     }
017     octree->build();
018 }
019
020 inline bool BVH::Extents::intersect(
021     const float *precomputedNumerator, const float *precomputeDenominator,
022     float &tNear, float &tFar, uint8_t &planeIndex)
023 {
024     __sync_fetch_and_add(&numRayVolumeTests, 1);
025     for (uint8_t i = 0; i < kNumPlaneSetNormals; ++i) {
026         float tn = (d[i][0] - precomputedNumerator[i]) / precomputeDenominator[i];
027         float tf = (d[i][1] - precomputedNumerator[i]) / precomputeDenominator[i];
028         if (precomputeDenominator[i] < 0) std::swap(tn, tf);
029         if (tn > tNear) tNear = tn, planeIndex = i;
030         if (tf < tFar) tFar = tf;
031         if (tNear > tFar) return false; // test for an early stop
032     }
033
034     return true;
035 }
036
037 const Object* BVH::intersect(const Ray<float> &ray, IsectData &isectData) const
038 {
039     const Object *hitObject = NULL;
040     float precomputedNumerator[BVH::kNumPlaneSetNormals], precomputeDenominator[BVH::kN
041     for (uint8_t i = 0; i < kNumPlaneSetNormals; ++i) {
042         precomputedNumerator[i] = dot(planeSetNormals[i], ray.orig);
043         precomputeDenominator[i] = dot(planeSetNormals[i], ray.dir);
044     }

```

```

044     }
045     #if 0
046     float tClosest = ray.tmax;
047     for (uint32_t i = 0; i < rc->objects.size(); ++i) {
048         __sync_fetch_and_add(&numRayVolumeTests, 1);
049         float tNear = -kInfinity, tFar = kInfinity;
050         uint8_t planeIndex;
051         if (extents[i].intersect(precomputedNumerator, precomputeDenominator, tNear, tFar,
052             IsectData isectDataCurrent;
053             if (rc->objects[i]->intersect(ray, isectDataCurrent)) {
054                 if (isectDataCurrent.t < tClosest && isectDataCurrent.t > ray.tmin) {
055                     isectData = isectDataCurrent;
056                     hitObject = rc->objects[i];
057                     tClosest = isectDataCurrent.t;
058                 }
059             }
060         }
061     }
062     #else
063     uint8_t planeIndex = 0;
064     float tNear = 0, tFar = ray.tmax;
065     if (!octree->root->extents.intersect(precomputedNumerator, precomputeDenominator, tNear, tFar,
066         || tFar < 0 || tNear > ray.tmax)
067         return NULL;
068     float tMin = tFar;
069     std::priority_queue<BVH::Octree::QueueElement> queue;
070     queue.push(BVH::Octree::QueueElement(octree->root, 0));
071     while(!queue.empty() && queue.top().t < tMin) {
072         const OctreeNode *node = queue.top().node;
073         queue.pop();
074         if (node->isLeaf) {
075             for (uint32_t i = 0; i < node->data.size(); ++i) {
076                 IsectData isectDataCurrent;
077                 if (node->data[i]->object->intersect(ray, isectDataCurrent)) {
078                     if (isectDataCurrent.t < tMin) {
079                         tMin = isectDataCurrent.t;
080                         hitObject = node->data[i]->object;
081                         isectData = isectDataCurrent;
082                     }
083                 }
084             }
085         }
086         else {
087             for (uint8_t i = 0; i < 8; ++i) {
088                 if (node->child[i] != NULL) {
089                     float tNearChild = 0, tFarChild = tFar;
090                     if (node->child[i]->extents.intersect(precomputedNumerator, precomputeDenominator, tNearChild, tFarChild, planeIndex)) {
091                         float t = (tNearChild < 0 && tFarChild >= 0) ? tFarChild : tNearChild;
092                         queue.push(BVH::Octree::QueueElement(node->child[i], t));
093                     }
094                 }
095             }
096         }
097     }
098     return hitObject;
099 }

```

```
096 |         }  
097 |     }  
098 | #endif  
099 |  
100 |     return hitObject;  
101 | }  
102 |
```

Normally, the objects which we should insert in the BVH are the **triangles** not the meshes. In order to complete the basic section as quickly as possible, we will explain how this can be done in a future version of this lesson. In the next chapter though, we will show how to insert the triangles in a grid acceleration structure. As an exercise, you can try to modify the code of the BVH found in this lesson to support the insertion of triangles, by taking example on the grid implementation.

Want to help?



## What's Next?

Kay and Kajiya's technique gives excellent results with the teapot scene. Generally BVH method perform well and are used sometimes in production renderers as the choice of supported accelerated structure (usually in combination with another structure for a reason that will become clear at the end of the next chapter).

The idea is to break these models somehow into smaller pieces which are inserted in the voxels a simple 3D grid. As we will show in the next chapter, ray tracing grids is simple and fast.

[← Previous Chapter](#)

Chapter 4 of 7

[Next Chapter →](#)